

Artificial Intelligence

Lecture 3

Supta Richard Philip

Department of CSE
City University

City University, December 2018



Table of Contents

1 Search Problem

2 References




Table of Contents

1 Search Problem

2 References



What is in State Space?



- A world state includes every details of the environment.
- A search state includes only details needed for planning.

Problem: Pathing

States: (x,y) locations

Actions: NSEW moves

Successor: update location

Goal: is (x,y) End?

Problem: Eat-all-dots

States: (x,y) , dot booleans


Actions: NSEW moves

Successor: update location and dot boolean

Goal: dots all false?



State Space Sizes?



World State: ?

$$120 * (2^{30}) * (12^2) * 4$$

States for Pathing: 120

$$\text{States for eat-all-dots: } 120 * (2^{30})$$

Pacman positions: $10 \times 12 = 120$

Pacman facing: up, down, left, right


Food Count: 30

Ghost positions: 12



State Space Graphs


- Each node is a state
- The successor function is represented by arcs
- Edges may be labeled with costs
- We can rarely build this graph in memory (so we don't)



Search Trees


A search tree:

- Start state at the root node
- Children correspond to successors
- Nodes contain states, correspond to PLANS to those states
- Edges are labeled with actions and costs
- For most problems, we can never actually build the whole tree




Example 1: Search Tree from state space graph

State Graph:




What is the search tree?

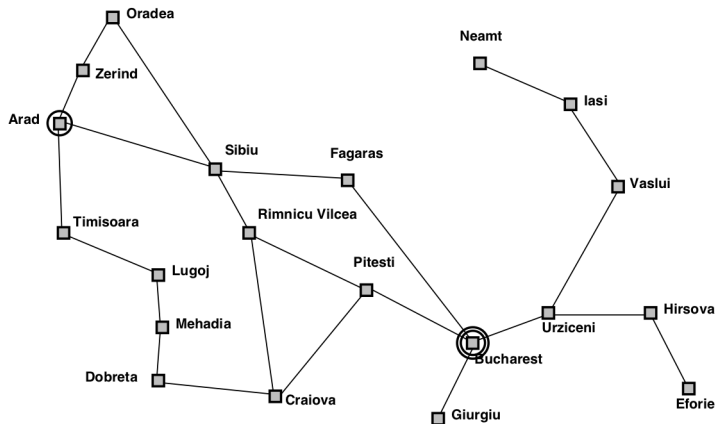
Path:



Search Tree:



Example 2: Search Tree from State Space Graph




Example 2: Search Tree from State Space Graph


(a) The initial state

Arad

(b) After expanding Arad



(c) After expanding Sibiu



General Search Algorithm

```
function GENERAL-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

```
function GENERAL-SEARCH(problem, QUEUEING-FN) returns a solution, or failure
  nodes  $\leftarrow$  MAKE-QUEUE(MAKE-NODE(INITIAL-STATE[problem]))
  loop do
    if nodes is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(nodes)
    if GOAL-TEST[problem] applied to STATE(node) succeeds then return node
    nodes  $\leftarrow$  QUEUEING-FN(nodes, EXPAND(node, OPERATORS[problem]))
  end
```



Search strategies

A strategy is defined by picking the *order of node expansion*.
Strategies are evaluated along the following dimensions:

- **completeness:** is the strategy guaranteed to find a solution when there is one?
- **time complexity:** how long does it take to find a solution?
- **space complexity:** how much memory is required to perform the search?
- **optimality:** does the search strategy find the highest quality solution when there are multiple solutions?

Time and space complexity are measured in terms of
 b —maximum branching factor of the search tree
 d —depth of the least-cost solution
 m —maximum depth of the state space (may be ∞)



Uninformed Search

Uninformed strategies use only the information available in the problem definition.

- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search





Depth-first search

Strategy: expand deepest node first

Implementation: Fringe is a LIFO queue (a stack)

Expansion order: (d,b,a,c,a,e,h,p,q,q,r,f,c,a,G)



Properties of Depth-first search

Complete: No. fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path

⇒ complete in finite spaces

Optimal: No

Time: $O(b^m)$: terrible if m is much larger than d

but if solutions are dense, may be much faster than breadth-first

Space: $O(bm)$, i.e., linear space!



Depth-limited search

depth-first search with depth limit /

Implementation: Nodes at depth / have no successors





Breadth-first search


Strategy: expand shallowest node first

Implementation: Fringe is a FIFO queue

Expansion order: (S,d,e,p,b,c,e,h,r,q,a,a,h,r,p,q,f,p,q,f,q,c,G)



Breadth-first search



Properties of Breadth-first search

Complete: Guaranteed to find a solution if one exists?

Yes (if b is finite)

Optimal: Guaranteed to find the least cost path?

Yes (if cost = 1 per step); not optimal in general

Time: $1 + b + b^2 + b^3 + \dots + b^d = O(b^d)$, i.e., exponential in d

Space: $O(b^d)$ (keeps every node in memory)

Space is the big problem; can easily generate nodes at 1MB/sec
so 24hrs = 86GB.






Table of Contents

1 Search Problem

2 References



References

-  Stuart Russell and Peter Norvig. 2009. Artificial Intelligence: A Modern Approach (3rd ed.). Prentice Hall Press, Upper Saddle River, NJ, USA.
-  <http://www.massey.ac.nz/~mjjohnso/notes/59302/all.html>